


# Concurrency Control

# Concurrency Control



# Transactions

A transaction is an execution of a program having the following properties:

## Atomicity

A transaction either happens or doesn't.

A transaction either completes and the results become visible or no results are visible.

## Consistency

Transactions preserve correctness of database.

## Isolation

Each transaction is unaware of other transactions executing concurrently

## Durability

The results of a completed transaction are permanently installed within D.B.

**ACID** properties

A transaction has exactly one of two possible outcomes:

a) Commit

Program execution completes and the results become permanent in database.

b) Abort

Program execution was not successful. “Results” are not installed into database.

- Transactions may abort due to hardware failure, system error, bad input data, or as a means to ensure consistency.

## Example:

T1: Read(A)  
A  $\leftarrow$  A+100  
Write(A)  
Read(B)  
B  $\leftarrow$  B+100  
Write(B)

T2: Read(A)  
A  $\leftarrow$  A $\times$ 2  
Write(A)  
Read(B)  
B  $\leftarrow$  B $\times$ 2  
Write(B)

Constraint: A=B

# Main idea of concurrency control

1. An execution without any interleaving is OK  
serial: T1 ; T2 or T2 ; T1
2. If an execution has the same effect as a serial execution then it is also acceptable  
serializable

Main goal of concurrency control:  
guarantee serializability

# Serial Schedule A ("good" by definition)

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			125
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		250
		250	250



# Serial Schedule B (equally "good")

		A	B
T1	T2	25	25
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	50	
	Read(B); $B \leftarrow B \times 2$ ;		50
	Write(B);		
Read(A); $A \leftarrow A + 100$			
Write(A);		150	
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			150
		150	150

# Interleaved Schedule C (good because it is equivalent to A)

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A+100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
Read(B); $B \leftarrow B+100$ ;			125
Write(B);			
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		250
		250	250

A and C are equivalent because if they start from same initial values they end up with same results

# Interleaved Schedule D (bad!)

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A+100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$ ;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$ ;		
	Write(B);		50
Read(B); $B \leftarrow B+100$ ;			
Write(B);			150
		250	150

Same as Schedule D  
but with new T2'

## Schedule E (good by "accident")

T1	T2'	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 1$ ;		
	Write(A);	125	
	Read(B); $B \leftarrow B \times 1$ ;		
	Write(B);		25
Read(B); $B \leftarrow B + 100$ ;			
Write(B);			125
		125	125

The accident being the particular semantics

- Want schedules that are “good”, I.e., equivalent to serial regardless of
  - initial state and
  - transaction semantics
- Only look at order of read and writes

## What we say to dogs

Okay, Ginger! I've had it!  
You stay out of the garbage!  
Understand, Ginger? Stay out  
of the garbage, or else!



## What they hear

blah blah GINGER blah  
blah blah blah blah blah  
blah blah GINGER blah  
blah blah blah blah...



What we say to a database

“delete all movies not  
directed by Berto”

---

read, read, write, read, read....

What the database hears

Example of a read/write schedule:

SC=r1(A)w1(A)r2(A)w2(A)r1(B)w1(B)r2(B)w2(B)

# Definition

$S_1, S_2$  are conflict equivalent schedules if  $S_1$  can be transformed into  $S_2$  by a series of swaps of adjacent non-conflicting actions.

Non-conflicting actions:

- actions on different data
- read/read on the same data




## Definition

A schedule is conflict serializable if it is conflict equivalent to some serial schedule.

Example:

$$SC = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$$



$$SC' = r_1(A)w_1(A) \ r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$$

$T_1$

$T_2$


However, for SD:

$SD = r_1(A)w_1(A)r_2(A)w_2(A) \ r_2(B)w_2(B)r_1(B)w_1(B)$

The diagram illustrates dependencies between two transactions,  $T_1$  and  $T_2$ . The schedule is  $SD = r_1(A)w_1(A)r_2(A)w_2(A) \ r_2(B)w_2(B)r_1(B)w_1(B)$ . Brackets group the operations for each transaction:  $\{r_1(A), w_1(A)\}$  for  $T_1$  and  $\{r_2(A), w_2(A), r_2(B), w_2(B)\}$  for  $T_2$ . Arrows with red X marks indicate dependencies that are not serializable: one arrow from  $r_1(A)$  to  $r_2(A)$  and another from  $w_2(B)$  to  $w_1(B)$ .

- as a matter of fact,  
 $T_2$  must precede  $T_1$   
in any equivalent schedule,  
i.e.,  $T_2 \rightarrow T_1$
- And vice versa

- $T_2 \rightarrow T_1$
- Also,  $T_1 \rightarrow T_2$




⇒ SD cannot be rearranged  
into a serial schedule

⇒ SD is not “equivalent” to  
any serial schedule

⇒ SD is “bad”

## Returning to SC



- no cycles  $\Rightarrow$  SC is "equivalent" to a serial schedule  
(in this case  $T_1, T_2$ )

# Precedence graph $P(S)$ ( $S$ is schedule)

Nodes: transactions in  $S$


Arcs:  $T_i \rightarrow T_j$  whenever

- $p_i(A), q_j(A)$  are actions in  $S$
- $p_i(A) <_S q_j(A)$
- at least one of  $p_i, q_j$  is a write

$<_S$  : order of appearance in schedule

## Exercise:

- What is  $P(S)$  for  
 $S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D)$



- Is  $S$  serializable?

## Lemma

$S_1, S_2$  conflict equivalent  $\Rightarrow P(S_1) = P(S_2)$

Proof:

Assume  $P(S_1) \neq P(S_2)$

$\Rightarrow \exists T_i, T_j: T_i \rightarrow T_j$  in  $S_1$  and not in  $S_2$

$\Rightarrow S_1 = \dots p_i(A) \dots q_j(A) \dots$	$\left\{ \begin{array}{l} p_i, q_j \\ \text{conflict} \end{array} \right.$
$S_2 = \dots q_j(A) \dots p_i(A) \dots$	

$\Rightarrow S_1, S_2$  not conflict equivalent



## Lemma

$S_1, S_2$  conflict equivalent  $\Rightarrow P(S_1)=P(S_2)$

Is the converse true?

A: yes      B: no

Note:  $P(S_1)=P(S_2) \not\Rightarrow S_1, S_2$  conflict equivalent

Counter example:

$S_1 = w_1(A) \ r_2(A) \quad w_2(B) \ r_1(B)$

$S_2 = r_2(A) \ w_1(A) \quad r_1(B) \ w_2(B)$

# Theorem

$P(S_1)$  acyclic  $\iff S_1$  conflict serializable

( $\Leftarrow$ ) Assume  $S_1$  is conflict serializable

$\Rightarrow \exists S_s: S_s, S_1$  conflict equivalent

$\Rightarrow P(S_s) = P(S_1)$

$\Rightarrow P(S_1)$  acyclic since  $P(S_s)$  is acyclic

# Theorem

$P(S_1)$  acyclic  $\iff S_1$  conflict serializable


$(\implies)$  Assume  $P(S_1)$  is acyclic

Transform  $S_1$  as follows:

(1) Take  $T_1$  to be transaction with no incoming arcs


(2) Move all  $T_1$  actions to the front

$S_1 = \dots\dots q_j(A)\dots\dots p_1(A)\dots\dots$



(3) we now have  $S_1 = \langle T_1 \text{ actions} \rangle \langle \dots \text{rest} \dots \rangle$

(4) repeat above steps to serialize rest!




# How to enforce serializable schedules?

*Option 1:* run system, recording  $P(S)$ ;  
check for  $P(S)$  cycles and  
declare if execution was good;  
or abort transactions as soon  
as they generate a cycle

# How to enforce serializable schedules?

*Option 2:* prevent P(S) cycles from occurring




# A locking protocol

Two new actions:

lock (exclusive):  $li(A)$

unlock:  $ui(A)$



# Rule #1: Well-formed transactions

$T_i: \dots l_i(A) \dots p_i(A) \dots u_i(A) \dots$



## Rule #2    Legal scheduler

$S = \dots \dots \dots l_i(A) \dots \dots \dots u_i(A) \dots \dots \dots$

$\longleftrightarrow$

no  $l_j(A)$

## Exercise:

**A:** S1 correct

**B:** S2 correct

**C:** S3 correct

- What schedules are correct?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

## Exercise:

- What schedules are correct?

S1 =  $l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$   
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

S2 =  $l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

S3 =  $l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$   
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$


Just having locks is not enough!

Example: this allows Schedule F (bad)


T1	T2
$l_1(A); \text{Read}(A)$ $A \leftarrow A + 100; \text{Write}(A); u_1(A)$	$l_2(A); \text{Read}(A)$ $A \leftarrow A \times 2; \text{Write}(A); u_2(A)$
$l_1(B); \text{Read}(B)$ $B \leftarrow B + 100; \text{Write}(B); u_1(B)$	$l_2(B); \text{Read}(B)$ $B \leftarrow B \times 2; \text{Write}(B); u_2(B)$

# Two phase locking (2PL)


$T_i = \dots\dots li(A) \dots\dots ui(A) \dots\dots$



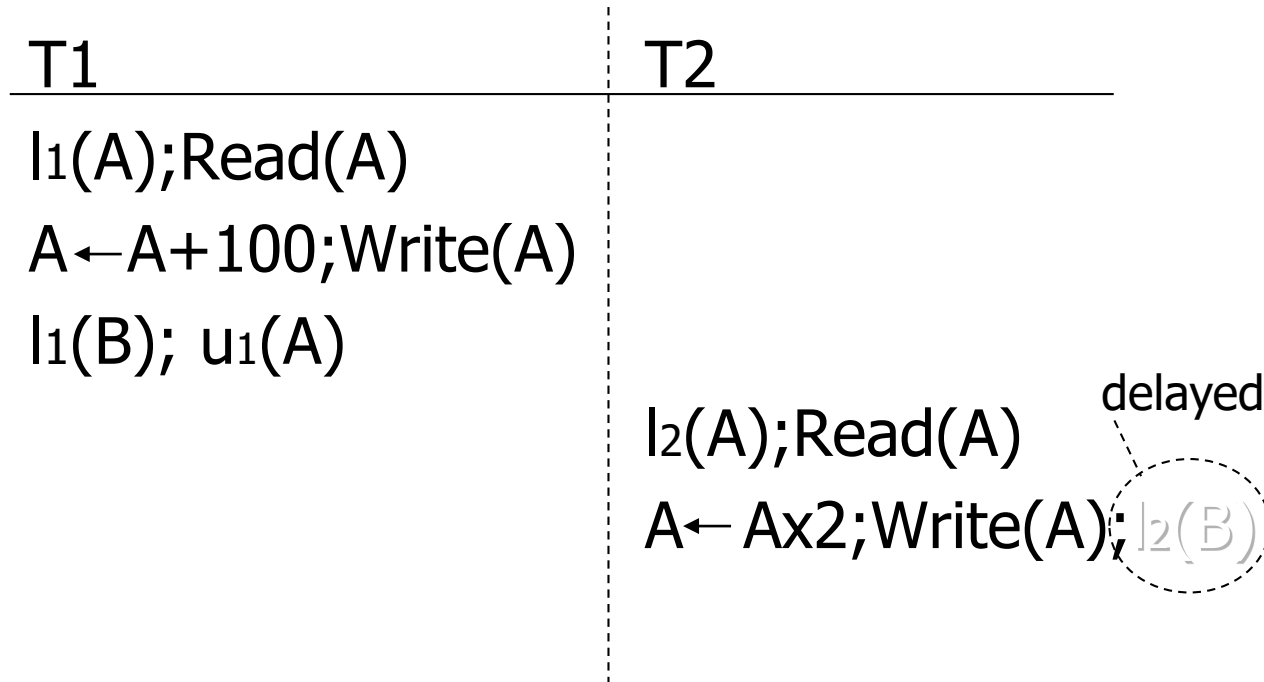
no unlocks



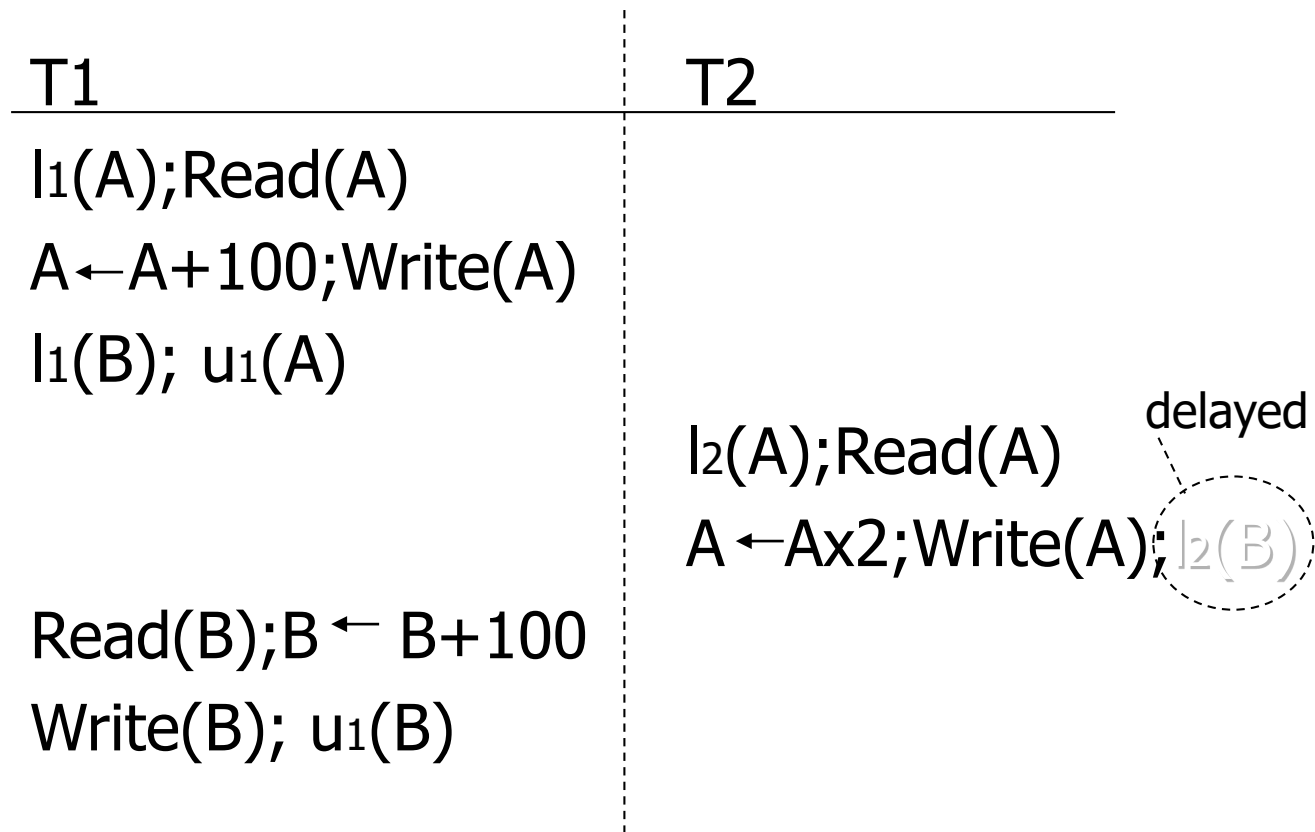
no locks



## 2PL prevents Schedule F:



# 2PL prevents Schedule F:






# 2PL prevents Schedule F:

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	
	$l_2(B); u_2(A); \text{Read}(B)$
	$B \leftarrow B \times 2; \text{Write}(B); u_2(B);$

## 2PL may however result in deadlock:



Usual fix: deadlocked transactions are aborted  
and rolled back

Next step:

Show that 2 Phase Locking  $\Rightarrow$  conflict-serializable schedules

Theorem 2PL  $\Rightarrow$  conflict serializable  
schedule

To help in proof:

Definition       $\text{Shrink}(Ti) = \text{SH}(Ti) =$   
first unlock action of  $Ti$

## Lemma

$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

### Proof of lemma:

$T_i \rightarrow T_j$  means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$



By 2PL:  $SH(T_i)$   $SH(T_j)$

$$\text{So, } SH(T_i) <_S SH(T_j)$$

Theorem 2PL  $\Rightarrow$  conflict serializable  
schedule

Proof:

(1) Assume  $P(S)$  has cycle

$$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$$

(2) By lemma:  $SH(T_1) < SH(T_2) < \dots < SH(T_1)$

(3) Impossible, so  $P(S)$  acyclic

(4)  $\Rightarrow S$  is conflict serializable


# How does locking work in practice?

- Every system is different

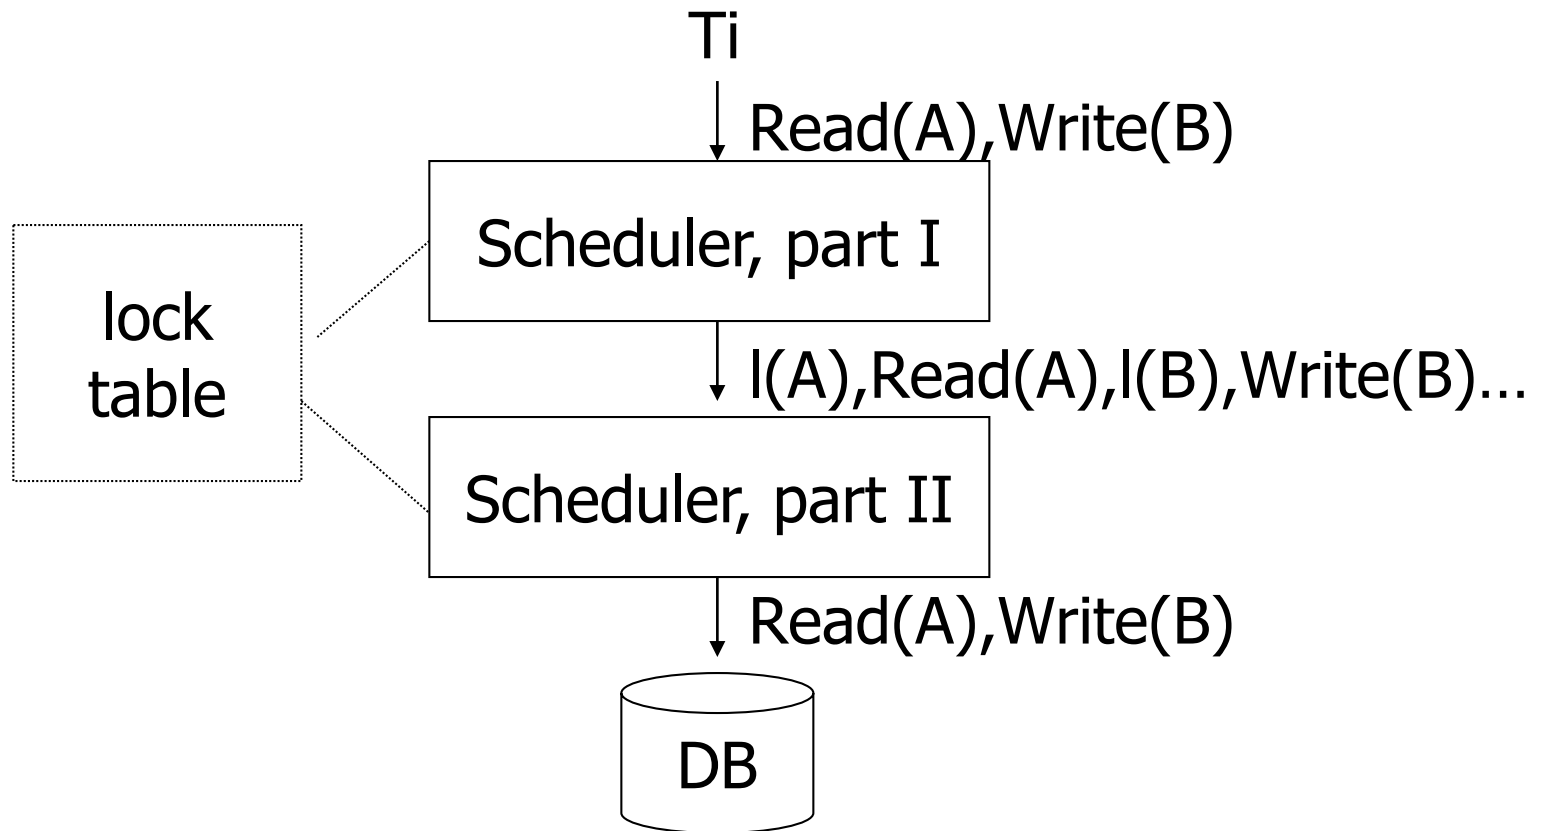
But here is one (simplified) way ...

## Sample Locking System:


- (1) Don't trust transactions to request/release locks
- (2) Hold all locks until transaction commits









# What are the objects we lock?



DB



DB



DB

?

- Locking works in any case, but should we choose small or large objects?
- If we lock large objects (e.g., Relations)
  - Need few locks
  - Low concurrency
- If we lock small objects (e.g., tuples, fields)
  - Need more locks
  - More concurrency

- Beyond this simple 2PL protocol, it is all a matter of improving performance and allowing more concurrency....
  - Shared locks
  - Multiple granularity
  - Inserts, deletes
  - Other types of C.C. mechanisms
  - Weaker guarantees (e.g. in the cloud)