

How expressive is SQL?

Full programming languages can express
all computable functions (C, Java, etc)

Can SQL express all computable queries?

A: YES B: NO

How expressive is SQL?

flight	from	to
	SD	LA
	SD	ORD
	LA	NY
	

Can SQL express the following query:
“Is there a way to get from City1 to City2” ?

A: YES B: NO

Limitation of basic SQL

flight	from	to
	SD	LA
	SD	ORD
	LA	NY
	

“Is there a way to get from City1 to City2” ?

Easier:

“Is there a way to get from City1 to City2 by a direct flight?”

City1 —————> City2

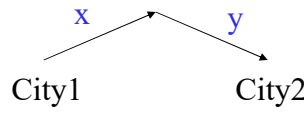
Select * from flight
where from = 'City1' and to = 'City2'

Easier:

“Is there a way to get from City1 to City2 **with at most one stopover?**”

City1 \longrightarrow City2 `select * from flight
where from = 'City1' and to = 'City2'`

OR

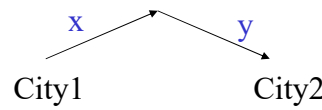
 `select x.from, y.to
from flight x, flight y
where x.from = 'City1' and
x.to = y.from and y.to = 'City2'`

Easier:

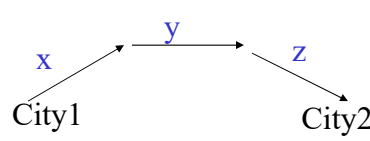
“Is there a way to get from City1 to City2 **with at most two stopovers?**”

City1 \longrightarrow City2 `select * from flight
where from = 'City1' and to = 'City2'`


OR

 `select x.from, y.to
from flight x, flight y
where x.from = 'City1' and
x.to = y.from and y.to = 'City2'`

OR

 `select x.from, z.to
from flight x, flight y, flight z
where x.from = 'City1' and x.to = y.from
and y.to = z.from and z.to = 'City2'`

“Is there a way to get from City1 to City2 **with at most k stopovers?**”



Need $k+1$ tuple variables!

“Is there a way to get from City1 to City2
with any number of stopovers?”

Cannot do in basic SQL!


Similar examples

- Parts-components relation:
“**find all subparts of some given part A**”
- Parent/child relation
“**find all of John’s descendants**”

More general: transitive closure of a graph

G	A	B
	a	b
	c	a
	b	c
	

Find the pairs of nodes $\langle x, y \rangle$ that are connected by some directed path



	a	b
	a	d
	a	e
	b	d
	b	e
	

Computing transitive closure T of G

“Find the pairs of nodes $\langle a, b \rangle$ that are connected in G”

Same as

“find pairs of nodes $\langle a, b \rangle$ at distance 1” UNION



“find pairs of nodes $\langle a, b \rangle$ at distance at most 2” UNION


.....

“find pairs of nodes $\langle a, b \rangle$ at distance at most k” UNION


.....

When to stop? At some point, no new nodes are added.
Distance cannot be larger than total number of nodes in G.





Distance ≤ 2




Distance ≤ 3

Denote by T_k the pairs of nodes at distance at most k

T_1 : “find pairs of nodes $\langle a, b \rangle$ at distance 1”

select * from G


T_k : “find the pairs of nodes $\langle a, b \rangle$ at distance at most k ”




(select * from T_{k-1})

union


(select x.A, y.B from G x, T_{k-1} y
where x.B = y.A)



T_1



T_2



T_3

One solution: extend SQL with recursion (not part of the standard)

```
create recursive view T as  
(select * from G)  
union  
(select x.A, y.B  
from G x, T y  
where x.B = y.A)
```

Semantics:

1. Start with empty T
2. While T changes
{evaluate view with current T;
union result with T }

Note: this must terminate, since
there are finitely many tuples
one can add to T (if no new values
are created)

Alternative

```
with recursive T as  
(select * from G)  
union  
(select x.A, y.B  
from G x, T y  
where x.B = y.A)  
select * from T ;
```

Semantics:

1. Start with empty T
2. While T changes
{evaluate view with current T;
union result with T }

Note: this must terminate, since
there are finitely many tuples
one can add to T (if no new values
are created)

Another example

frequents	drinker	bar
-----------	---------	-----

Friends: drinkers who frequent the same bar
Find transitive closure of **Friends**

```
create recursive view T as
select f1.drinker as drinker1, f2.drinker as drinker2
from frequents f1, frequents f2
where f1.bar = f2.bar
union
select t1.drinker1, f2.drinker as drinker 2
from T t1, frequents f1, frequents f2
where t1.drinker2 = f1.drinker and f1.bar = f2.bar
```

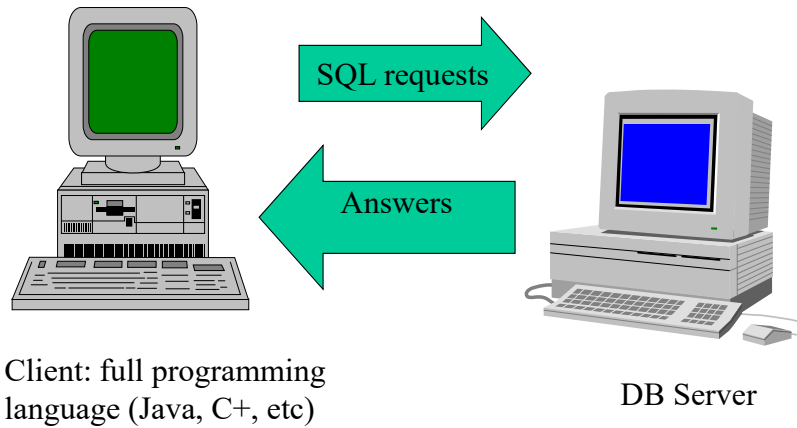
Problematic example

```
create recursive view T as
select A, 0 as N from R
union
select A, N+1 as N
from T
```

R	A
---	---

- never terminates
- some systems forbid arithmetic or aggregate functions in selects in recursive definitions

Powerful way to overcome limitations
embedded SQL




23


Transitive closure in embedded SQL

<pseudo-code>


```
T := G
Δ := G
while Δ ≠ Φ do
  { Told = T
    T := (select * from T)
        union
        (select x.A, y.B from G x, T y
         where x.B = y.A)
    Δ := T - Told }
Output T
```




T_1 and Δ_1



T_2 and Δ_2



T_3 and Δ_3



$T_4 = T_3$ and $\Delta_4 = \Phi$
 stop!

Another way: Embedded SQL

<pseudo-code>


```
T := G
Δ := G
while Δ ≠ Φ do
  { Told = T
    T := (select * from T)
         union
         (select x.A, y.B from G x, T y
          where x.B = y.A)
    Δ := T - Told }
Output T
```

Converges in **diameter(G)** iterations
(maximum distance between two nodes in G)


Optimization (“semi-naïve” evaluation):
use at least one new tuple (from **Δ**) every time

<pseudo-code>

```
T := G
Δ := G
while Δ ≠ Φ do
  { Told = T
    T := (select * from T)
         union
         (select x.A, y.B from G x, Δ y
          where x.B = y.A)
    Δ := T - Told }
Output T
```




T_1 and Δ_1




T_2 and Δ_2

No longer recompute $\langle c, b \rangle$ but recompute $\langle c, d \rangle$



T_3 and Δ_3

No longer recompute $\langle a, d \rangle$ but recompute $\langle c, e \rangle$



$T_4 = T_3$ and $\Delta_4 = \Phi$

stop!

Faster convergence (double recursion):

```
T := G
Δ := G
while Δ ≠ Φ do
  { Told = T
    T := (select * from T)
        union
        (select x.A, y.B from T x, T y
         where x.B = y.A)
    Δ := T - Told }
```


Output T

Converges in $\log(\text{diameter}(G))$ iterations


Example (focus on computing $\langle a_0, a_8 \rangle$)

$a_0 \longrightarrow a_1 \longrightarrow a_2 \longrightarrow a_3 \longrightarrow a_4 \longrightarrow a_5 \longrightarrow a_6 \longrightarrow a_7 \longrightarrow a_8$


Example (focus on computing $\langle a_0, a_8 \rangle$)




Example (focus on computing $\langle a_0, a_8 \rangle$)




Example (focus on computing $\langle a_0, a_8 \rangle$)




Compare to linear recursion (first program):




Compare to linear recursion (first program):




Compare to linear recursion (first program):




Compare to linear recursion (first program):




Compare to linear recursion (first program):




Compare to linear recursion (first program):




Compare to linear recursion (first program):



Compare to linear recursion (first program):



Compare to linear recursion (first program):



Optimization (“semi-naïve” evaluation):
again, use at least one new tuple every time

```
T := G
Δ := G
while Δ ≠ Φ do
  { Told = T
    T := (select * from T)
         union
         (select x.A, y.B from Δ x, T y
          where x.B = y.A)
         union
         (select x.A, y.B from T x, Δ y
          where x.B = y.A) }
  Δ := T - Told
Output T
```