

The Relational Model

- database consists of several **tables (relations)**
- columns in each table are named by **attributes**
- each attribute has an associated **domain**
(set of allowed values)
- data in each table consists of a set of **rows (tuples)** providing values for the attributes

Example

STUDENT	Name	SSN	HomePhone	Address	OfficePhone	Age	GPA
	Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	null	19	3.21
	Katherine Ashly	381-62-1245	375-4409	125 Kirby Road	null	18	2.89
	Dick Davidson	422-11-2320	null	3452 Elgin Road	749-1253	25	3.53
	Charles Cooper	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
	Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	null	19	3.25

Relation Schema

“type declaration”

- Relation name
- Set of attributes
- Domain of each attribute
- Integrity constraints

Example

CUSTOMER (Cust-id, Cust-name, Address, Phone#)
integer char strings 7-digits

Attribute Types

- Each attribute of a relation has a name
- The set of allowed values for each attribute is called the **domain** of the attribute
- Attribute values are (normally) required to be **atomic**; that is, indivisible
- Sometimes, the special value **null** is considered a member of every domain

Relation Instance

An **instance** of a relation schema is the current content of the relation: a **set of rows** (tuples) over the attributes, with values from the attribute domains

<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
<i>Jones</i>	Main	Harrison
<i>Smith</i>	North	Rye
<i>Curry</i>	North	Rye
<i>Lindsay</i>	Park	Pittsfield

More on tuples

Notation:

- We refer to **component values** of a tuple t by $t(A_i) = v_i$ (the value of attribute A_i for tuple t).
also called **coordinates**

Example

	<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>
t →	Jones Smith Curry Lindsay	Main North North Park	Harrison Rye Rye Pittsfield

$t = \langle \text{Smith, North, Rye} \rangle$

$t(\text{customer_name}) = \text{Smith}$

$t(\text{customer_street}) = \text{North}$

$t(\text{customer_city}) = \text{Rye}$

attributes and tuple
values are generally
assumed to be ordered

Relations are Unordered Sets

The tuples are *not* considered to be ordered, even though they appear to be so in the displayed tabular form.

<i>account_number</i>	<i>branch_name</i>	<i>balance</i>
A-101	Downtown	500
A-215	Mianus	700
A-102	Perryridge	400
A-305	Round Hill	350
A-201	Brighton	900
A-222	Redwood	700
A-217	Brighton	750

Alternative: multiset (bag) semantics

R	A	B
	1	1
	1	1
	0	1

R	A	B
	1	1
	1	1
	1	1
	0	1
	0	1

- same under set semantics
- different under multiset semantics
(takes into account number of occurrences)

Database

- A database consists of one or several relations
- Information about an application is usually broken up into parts, with each relation storing one part of the information
 - account* : stores information about accounts
 - depositor* : stores information about which customer owns which account
 - customer* : stores information about customers
- Storing all information as a single relation such as
bank (*account_number*, *balance*, *customer_name*, ..)
 is possible but not desirable:
 results in repetition of information and the need for null values

Relational Integrity Constraints

- Constraints are *conditions* that must hold on *all* valid relation instances of a database
- Some common types of constraints:
 1. **Key** constraints
 2. **Entity integrity** constraints
 3. **Referential integrity** constraints

Key Constraints

- **Superkey** of R: A set of attributes SK of R such that no two tuples in any valid relation instance $r(R)$ will have the same value for SK. That is, for all distinct tuples t_1 and t_2 in $r(R)$, $t_1(SK) \neq t_2(SK)$.
- **Key** of R: A "minimal" superkey; that is, a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey.

Example: The CAR relation schema:
 CAR(State, Reg#, SerialNo, Make, Model, Year)
 has two keys Key1 = {State, Reg#}, Key2 = {SerialNo}.
 {SerialNo, Make} is a superkey but *not* a key.
- If a relation has *several* **candidate keys**, one is chosen to be the **primary key**.

Key Constraints

CAR	<u>LicenseNumber</u>	EngineSerialNumber	Make	Model	Year
	Texas ABC-739	A69352	Ford	Mustang	96
	Florida TVP-347	B43696	Oldsmobile	Cutlass	99
	New York MPO-22	X83554	Oldsmobile	Delta	95
	California 432-TFY	C43742	Mercedes	190-D	93
	California RSK-629	Y82935	Toyota	Camry	98
	Texas RSK-629	U028365	Jaguar	XJS	98

The primary key attributes are *underlined*.

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

ESSN	PNO	HOURS
------	-----	-------

DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
John			Smith	123456789	1965-01-09	731 Forden, Houston, TX	M	30000	333445555	5
Franklin			Wong	333445555	1965-12-08	638 Voss, Houston, TX	M	40000	888885555	5
Alicia			Zelaya	999987777	1968-01-19	3321 Casale, Spring, TX	F	25000	987654321	4
Jennifer			Wallace	987654321	1941-05-20	291 Berry, Bellaire, TX	F	43000	888885555	4
Ramesh			Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	36000	333445555	5
Joyce			English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad			Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James			Borg	888885555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1

WORKS_ON	ESSN	PNO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40.0
	453453453	1	20.0
	453453453	2	20.0
	333445555	2	10.0
	333445555	3	10.0
	333445555	10	10.0
	333445555	20	10.0
	999987777	30	30.0
	999987777	10	10.0
	987987987	10	35.0
	987987987	30	5.0
	987654321	30	20.0
	987654321	20	15.0
	888885555	20	null

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATIONSHIP
	333445555	Alice	F	1988-04-05	DAUGHTER
	333445555	Theodore	M	1983-10-25	SON
	333445555	Joy	F	1958-05-03	SPOUSE
	987654321	Abner	M	1942-02-28	SPOUSE
	123456789	Michael	M	1988-01-04	SON
	123456789	Alice	F	1988-12-30	DAUGHTER
	123456789	Elizabeth	F	1967-05-05	SPOUSE

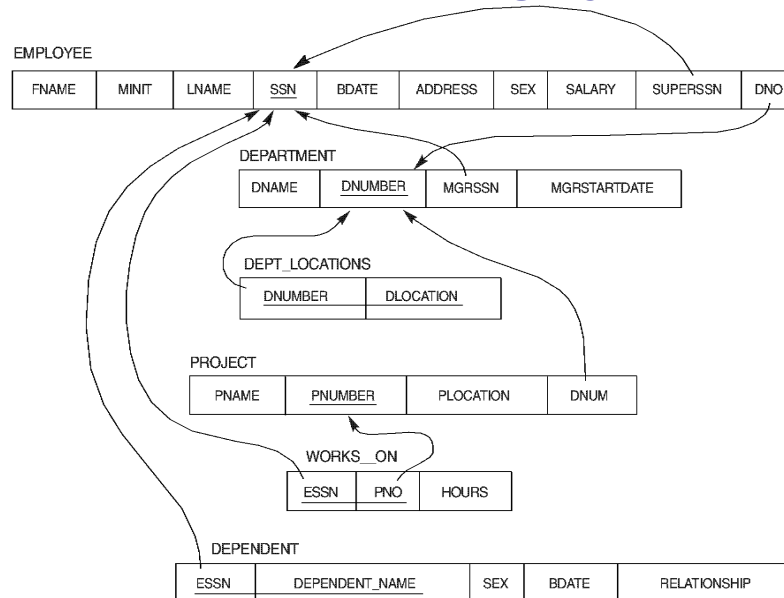
Entity Integrity

- The *primary key attributes* PK of each relation schema R in S cannot have null values in any tuple. This is because PK values are used to *identify* the individual tuples.

$t(A) \neq \text{null}$ for every tuple t in
a valid instance of R, where A is in PK

Note: Other attributes of R may be similarly constrained to disallow null values, even though they are not members of the primary key.

Referential Integrity



Referential Integrity

- A constraint involving *two* relations of the database (the previous constraints involve a *single* relation).
- Used to specify a *relationship* among tuples in two relations: the **referencing relation** and the **referenced relation**.
- Tuples in the *referencing relation* R_1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the *referenced relation* R_2 . A tuple t_1 in R_1 is said to **reference** a tuple t_2 in R_2 if $t_1(\text{FK}) = t_2(\text{PK})$.
- A referential integrity constraint can be displayed in a relational database schema as a directed arc from $R_1.\text{FK}$ to $R_2.\text{PK}$.

Referential Integrity Constraint

Statement of the constraint

The value in the foreign key column(s) FK of the **referencing relation** R_1 can be either (1) a value of a primary key PK in the **referenced relation** R_2 or (2) null.

Other Types of Constraints

- Semantic Integrity Constraints: based on application semantics and cannot be expressed by the model per se
 - e.g., “the max. no. of hours per employee for all projects he or she works on is 40 hrs per week”
 - A *constraint specification language* may have to be used to express these

SQL provides assertions and triggers

Update Operations on Relations

- INSERT a tuple.
- DELETE a tuple.
- MODIFY a tuple.

- Integrity constraints should not be violated by the update operations.
- Several update operations may have to be grouped together.

Update Operations on Relations

- In case of integrity violation, several actions can be taken:
 - Cancel the operation that causes the violation (REJECT option)
 - Perform the operation but inform the user of the violation
 - Trigger additional updates so the violation is corrected
 - Execute a user-specified error-correction routine

SQL

“Structured Query Language”

- Standard for relational db systems
 - History:
 - Developed at IBM in late 70s
 - First standard: SQL-86
 - Second standard: SQL-92
 - Third standard: SQL-99 or SQL3, well over 1000 pages!
 - Many more
- “The nice things about standards is that you have so many to choose from” -- Andres S. Tannenbaum

23

SQL Data Definition Language

Allows the specification of the database schema

- The name and attributes for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- The set of indices to be maintained for each relations.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk.

24

Some Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p,d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.

25

Create Table Command

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,
               (integrity-constraint1),
               ...,
               (integrity-constraintk))
```

- *r* is the name of the relation
- each *A_i* is an attribute name in the schema of relation *r*
- *D_i* is the domain of attribute *A_i*

- Example:

```
create table branch
  (branch_name char(15) not null,
   branch_city char(30),
   assets integer)
```

26

Create Table (cont.)

- Can use the CREATE TABLE command for specifying the primary key attributes, secondary keys, and referential integrity constraints (foreign keys).
- Key attributes can be specified via the PRIMARY KEY and UNIQUE keywords

```
CREATE TABLE DEPT
(  DNAME          VARCHAR(10)  NOT NULL,
   DNUMBER        INTEGER      NOT NULL,
   MGRSSN         CHAR(9),
   MGRSTARTDATE   CHAR(9),
   PRIMARY KEY    (DNUMBER),
   UNIQUE         (DNAME),
   FOREIGN KEY    (MGRSSN) REFERENCES EMP );
```

primary key declaration on an attribute automatically ensures **not null** in SQL-92 onwards, needs to be explicitly stated in SQL-89

27

The check clause

- **check** (*P*), where *P* is a predicate on attribute values

Declare *branch_name* as the primary key for *branch* and ensure that the values of *assets* are non-negative.

```
create table branch
(branch_name char(15),
 branch_city char(30),
 assets integer,
 primary key (branch_name),
 CHECK      (assets >= 0) )
```

28

Drop Table Command

- Used to remove a relation *and its definition*
- The relation can no longer be used in queries, updates, or any other commands since its description no longer exists
- Example:

DROP TABLE DEPENDENT;

29

Alter Table Command

- The **alter table** command is used to add attributes to an existing relation:
alter table r add A D
 where A is the name of the attribute to be added to relation r and D is the domain of A .
*All tuples in the relation are assigned **null** as the default value for the new attribute.*
- The **alter table** command can also be used to drop attributes of a relation:
alter table r drop A
 where A is the name of an attribute of relation r
Many databases do not support dropping of attributes

30

Alter Table (cont.)

- Since new attribute will have NULL values right after the **ALTER** command is executed, the NOT NULL constraint is *not allowed* for such an attribute
- Example:
**ALTER TABLE EMPLOYEE
ADD JOB VARCHAR(12);**
- The database users must still enter a value for the new attribute JOB for each EMPLOYEE tuple. This can be done using the UPDATE command.