

Null values in SQL

Relation name

Attributes

STUDENT	Name	SSN	HomePhone	Address	OfficePhone	Age	GPA
	Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	null	19	3.21
	Katherine Ashly	381-62-1245	375-4409	125 Kirby Road	null	18	2.89
	Dick Davidson	422-11-2320	null	3452 Elgin Road	749-1253	25	3.53
	Charles Cooper	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
	Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	null	19	3.25

Tuples

Null values in SQL

- testing if an attribute A is null: **A is null**, **A is not null**

Ex: find all employees with unknown phone number

```
select name from employee  
where phone is null
```

- arithmetic operations involving any null return null

if Salary is null, then $\text{Salary} + 1$ evaluates to null

- comparisons involving any null return **unknown**

new truth value

if Salary is null, then $\text{Salary} = 0$ evaluates to unknown

Null values in SQL

- Boolean operations must now handle 3 truth values:
true, false, unknown
- Boolean expressions involving **unknown** are evaluated using the following truth tables:

AND			
	true	unknown	unknown
	false	unknown	false
	unknown	unknown	unknown

NOT		
	unknown	unknown

OR			
	true	unknown	true
	false	unknown	unknown
	unknown	unknown	unknown

- WHERE clause conditions evaluating to **unknown** are treated as **false**

Examples

Movie	title	director	actor
	Tango	Berto	Brando
	Psycho	Hitch	Perkins
	Bambi	null	null

Select title
Where dir = 'Hitch'

title
Psycho

Select title
Where dir \neq 'Hitch'

title
Tango
Bambi

A: yes
B: no

Examples

Movie	title	director	actor
	Tango	Berto	Brando
	Psycho	Hitch	Perkins
	Bambi	null	null

Select title
Where dir = 'Hitch'

title
Psycho

Select title
Where dir \neq 'Hitch'

title
Tango
Bambi

A: yes
B: no

title
Tango

B

Examples

Movie	title	director	actor
	Tango	Berto	Brando
	Psycho	Hitch	Perkins
	Bambi	null	null

Select title
Where dir = 'null'

title
Bambi

A: yes
B: no

Examples

Movie	title	director	actor
	Tango	Berto	Brando
	Psycho	Hitch	Perkins
	Bambi	null	null

Select title
Where dir = 'null'

title
Bambi

A: yes
B: no

Select title
Where dir **is null**

title
Bambi

Null values in SQL

Anomalies of null semantics

if Salary is null, then:

- Salary * 0 evaluates to null
- Salary > 0 evaluates to unknown even if the domain is restricted to positive integers in the schema definition
- Consider the queries

select name from employee
where Salary <= 100 OR Salary > 100

and

select name from employee

Are these equivalent? A: yes B: no

Null values in SQL

Anomalies of null semantics

if Salary is null, then:

- $\text{Salary} * 0$ evaluates to null
- $\text{Salary} > 0$ evaluates to unknown even if the domain is restricted to positive integers in the schema definition
- Consider the queries

select name from employee
where Salary \leq 100 OR Salary $>$ 100

and

select name from employee

These are **not** equivalent if some salaries are null

Null Values and Aggregates

- Total all loan amounts

```
select sum (amount )  
from loan
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.

Suppose R has a single attribute A. Are these equivalent?

```
select count(*) from R
```

```
select count(A) from R
```

A: yes

B: no

Null Values and Group-By

- Null group-by attributes are treated like any other value

R	A	B
	2	3
	2	5
	Null	0
	Null	1
	Null	2

```
SELECT A, COUNT(B) AS C  
GROUP BY A
```

A	C
2	2
Null	3

SQL: Natural Join

Combines tuples from two tables by matching on common attributes

movie	title	director	actor
	Tango	Berto	Brando
	Sky	Berto	Winger
	Psycho	Hitchcock	Perkins

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Bambi
	Ken	Psycho

movie	natural join	schedule	title	director	actor	theater
			Tango	Berto	Brando	Hillcrest
			Tango	Berto	Brando	Paloma
			Psycho	Hitchcock	Perkins	Ken

SQL: Natural Join

Find the directors of movies showing in Hillcrest:

```
select director  
from movie natural join schedule  
where theater = 'Hillcrest'
```

Many variations of joins are available in SQL

SQL: creating nulls with Outer Joins

Idea: to avoid losing tuples in natural joins, pad with **null values**

P <outer join> Q

- **natural left outer join**: keep all tuples from left relation (P)
- **natural right outer join**: keep all tuples from right relation (Q)
- **natural full outer join**: keep all tuples from both relations

SQL: creating nulls with Outer Joins

Combines tuples from two tables by matching on common attributes

movie	title	director	actor
	Tango	Berto	Brando
	Sky	Berto	Winger
	Psycho	Hitchcock	Hopkins

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Bambi
	Ken	Psycho

movie	natural left outer join	title	director	actor	theater
schedule		Tango	Berto	Brando	Hillcrest
		Tango	Berto	Brando	Paloma
		Psycho	Hitchcock	Hopkins	Ken
		Sky	Berto	Winger	<i>null</i>

SQL: creating nulls with Outer Joins

Combines tuples from two tables by matching on common attributes

movie	title	director	actor
	Tango	Berto	Brando
	Sky	Berto	Winger
	Psycho	Hitchcock	Hopkins

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Bambi
	Ken	Psycho

movie	natural right outer join	title	director	actor	theater
schedule		Tango	Berto	Brando	Hillcrest
		Tango	Berto	Brando	Paloma
		Psycho	Hitchcock	Hopkins	Ken
		Bambi	<i>null</i>	<i>null</i>	Paloma

SQL: creating nulls with Outer Joins

Combines tuples from two tables by matching on common attributes

movie	title	director	actor
	Tango	Berto	Brando
	Sky	Berto	Winger
	Psycho	Hitchcock	Hopkins

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Bambi
	Ken	Psycho

movie	natural full outer join	title	director	actor	theater
schedule		Tango	Berto	Brando	Hillcrest
		Tango	Berto	Brando	Paloma
		Psycho	Hitchcock	Hopkins	Ken
		Bambi	<i>null</i>	<i>null</i>	Paloma
		Sky	Berto	Winger	<i>null</i>

Example: Find theaters showing only movies by Berto

```
select theater from schedule
where theater not in
  (select theater
   from schedule natural left outer join
     (select title, director from movie where director = 'Berto')
   where director is null)
```

```
select title, director from movie where director = 'Berto'
```

Movie	title	director	actor
	Tango	Berto	Brando
	Sky	Berto	Winger
	Psycho	Hitchcock	Hopkins

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Psycho

Example: Find theaters showing only movies by Berto

```
select theater from schedule
where theater not in
  (select theater
   from schedule natural left outer join
     (select title, director from movie where director = 'Berto')
   where director is null)
```

```
schedule natural left outer join (select title, director from movie
                                   where director = 'Berto')
```

title	director
Tango	Berto
Sky	Berto

schedule	theater	title
	Hillcrest	Tango
	Paloma	Tango
	Paloma	Psycho

theater	title	director
Hillcrest	Tango	Berto
Paloma	Tango	Berto
Paloma	Psycho	null

Summary of basic SQL Queries

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory.
- The clauses are specified in the following order:

SELECT <attribute list>
FROM <table list>
[WHERE <condition>]
[GROUP BY <grouping attribute(s)>]
[HAVING <group condition>]
[ORDER BY <attribute list>]

Summary of basic SQL Queries (cont.)

- The SELECT-clause lists the attributes or functions to be retrieved
- The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The WHERE-clause specifies the conditions for selection of tuples from the relations specified in the FROM-clause
- GROUP BY specifies grouping attributes
- HAVING specifies a condition for selection of groups
- ORDER BY specifies an order for displaying the result of a query
- A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause

SQL update language

Insertions

- inserting tuples
 - INSERT INTO *R*
VALUES (v_1, \dots, v_k);
- some values may be left NULL
- use results of queries for insertion
 - INSERT INTO *R* SELECT ...
FROM ... WHERE

```
INSERT INTO Movie  
VALUES ("Matchpoint", "Allen", "Allen");
```

```
INSERT INTO Movie(Title,Director)  
VALUES ("Matchpoint", "Allen");
```

```
INSERT INTO BertoMovie  
SELECT * FROM Movie  
WHERE Director = "Berto"
```

SQL update language:

Updates and Deletions

- **Deletion** basic form: delete every tuple that satisfies *<cond>*
 - DELETE FROM *R* WHERE *<cond>*
- **Update** basic form: update every tuple that satisfies *<cond>* in the way specified by the SET clause
 - UPDATE *R*
SET *A1*=*<expl>*,
...
Ak=*<expk>*
WHERE *<cond>*

Delete the movies that are not currently playing

```
DELETE FROM Movie  
WHERE Title NOT IN SELECT Title  
FROM Schedule
```

Change all “Berto” entries to “Bertolucci”

```
UPDATE Movie  
SET Director=“Bertolucci”  
WHERE Director=“Berto”
```

Increase all salaries in the Toys dept by 10%

```
UPDATE Employee  
SET Salary = 1.1 * Salary  
WHERE Dept = “Toys”
```

The “rich get richer” exercise:

Increase by 10% the salary of the employee with the highest salary

Example: delete all theaters showing more than one title

```
delete from schedule s
where exists (select * from schedule
              where theater = s.theater and title <> s.title)
```

Schedule	theater	title
	Hillcrest	Amour
	Hillcrest	0 dark 30
	Paloma	Django

Assume this semantics:
for each s in schedule
if where clause is satisfied
then delete s



Result after delete?

A: yes B: no

Schedule	theater	title
	Paloma	Django

Example: delete all theaters showing more than one title

```
delete from schedule s
where exists (select * from schedule
              where theater = s.theater and title <> s.title)
```

Schedule	theater	title
	Hillcrest	Amour
	Hillcrest	0 dark 30
	Paloma	Django

Correct semantics:

1. Find all theaters showing more than one title
2. Delete all theaters found in 1.



Result after delete?

Schedule	theater	title
	Paloma	Django

A: yes B: no