

# SQL Data Manipulation Language

- primarily **declarative** query language
- starting point: **relational calculus**  
aka first-order predicate logic
- many additions, bells and whistles...
- corresponding procedural language: **relational algebra**

will discuss relational calculus and algebra later

# Running example: Movie database

Movie	Title	Director	Actor

Schedule	Theater	Title

# Running example: Movie database

Movie	Title	Director	Actor
	Star Wars	Lucas	Ford
	Star Wars	Lucas	Fischer
	Mad Max	Miller	Hardy
	.....		

Schedule	Theater	Title
	Hillcrest	Star Wars
	Hillcrest	Mad Max
	Paloma	Rocky Horror
	.....	

*Find titles of currently playing movies*

```
SELECT Title  
FROM Schedule
```

*Find the titles of all movies by “Lucas”*

```
SELECT Title  
FROM Movie  
WHERE Director=“Lucas”
```

*Find the titles and the directors of all  
currently playing movies*

```
SELECT Movie.Title, Director  
FROM Movie, Schedule  
WHERE Movie.Title=Schedule.Title
```

# Basic form

```
SELECT  $a_1, \dots, a_n$   
FROM  $R_1, \dots, R_m$   
WHERE condition
```

WHERE clause is optional

# Informal semantics of basic form

```
SELECT  $a_1, \dots, a_n$   
FROM  $R_1, \dots, R_m$   
WHERE condition
```

```
for each tuple  $t_1$  in  $R_1$   
  for each tuple  $t_2$  in  $R_2$   
    .....  
    for each tuple  $t_m$  in  $R_m$   
  
      if  $\text{condition}(t_1, t_2, \dots, t_m)$  then output in answer  
        attributes  $a_1, \dots, a_n$  of  $t_1, \dots, t_m$ 
```

# Examples revisited

```
SELECT Title  
FROM Movie  
WHERE Director= "Lucas"
```

Semantics:

for each tuple  $m$  in Movie  
if  $m(\text{Director}) = \text{"Lucas"}$  then output  $m(\text{Title})$

# Examples revisited

```
SELECT Movie.Title, Director  
FROM Movie, Schedule  
WHERE Movie.Title=Schedule.Title
```

## Semantics:



```
for each tuple m in Movie  
  for each tuple s in Schedule  
    if m(title) = s(title) then output <m(Title),m(Director)>
```



# SQL Queries: Tuple variables

- Needed when using the same relation more than once in the FROM clause
- Example: **find actors who are also directors**

```
SELECT t.Actor
FROM Movie t, Movie s
WHERE t.Actor = s.Director
```

Movie	Title	Director	Actor
t			
s			

## Semantics:

for each t in Movie

for each s in Movie

if  $t(\text{Actor}) = s(\text{Director})$  then output  $t(\text{Actor})$

# Previous examples using tuple variables

```
SELECT Title  
FROM Movie  
WHERE Director= "Lucas"
```

```
SELECT m.Title  
FROM Movie m  
WHERE m.Director = "Lucas"
```

# Previous examples using tuple variables

```
SELECT Movie.Title, Director  
FROM Movie, Schedule  
WHERE Movie.Title=Schedule.Title
```

```
SELECT m.Title, m.Director  
FROM Movie m, Schedule s  
WHERE m.Title = s.Title
```

# SQL Queries: \* and LIKE

- Select all attributes using \*

*Retrieve all movie attributes of currently playing movies*

```
SELECT Movie.*  
FROM Movie, Schedule  
WHERE Movie.Title=Schedule.Title
```

- Pattern matching conditions
  - *<attr> LIKE <pattern>*

*Retrieve all movies where the title starts with “Ta”*

```
SELECT *  
FROM Movie  
WHERE Title LIKE “Ta%”
```

*Forgot if “Polanski” is spelled with “i” or “y”:*

```
SELECT *  
FROM Movie  
WHERE Director LIKE “Polansk_”
```

# SQL Queries: duplicate elimination

- *Default: answers to queries contain **duplicates***
- *Duplicate elimination must be explicitly requested*
  - SELECT DISTINCT ...  
FROM ... WHERE ...

```
SELECT Title  
FROM Movie
```

Title
Tango
Tango
Tango

```
SELECT DISTINCT Title  
FROM Movie
```

Title
Tango

# Ordering the Display of Tuples

- List all titles and actors of movies by Fellini, in alphabetical order of titles

```
select Title, Actor  
from   Movie  
where Director = 'Fellini'  
ORDER BY Title
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute;
  - ascending order is the default.
  - Example: **order by** *Title* **desc**

# Renaming attributes in result

Done using the **as** construct:

*Find titles of movies by Bertolucci, under attribute Berto-titles:*

```
select title as Berto-title  
from movie  
where director = 'Bertolucci'
```

# Aggregate Functions

These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values



## Aggregate Functions (Cont.)

-- Find the average account balance at the La Jolla branch.

```
select avg (balance)  
from account  
where branch_name = 'La Jolla'
```

-- Find the number of tuples in the *customer* relation.

```
select count (*)  
from customer
```

-- Find the number of depositors in the bank.

```
select count (distinct customer_name)  
from depositor
```

# Aggregate Functions (Cont.)

- Find the maximum salary, the minimum salary, and the average salary among all employees for the Company database

```
SELECT    MAX(SALARY),  
          MIN(SALARY), AVG(SALARY)  
FROM EMPLOYEE
```

Obs. Some SQL implementations *may not allow more than one aggregate function* in the SELECT-clause!

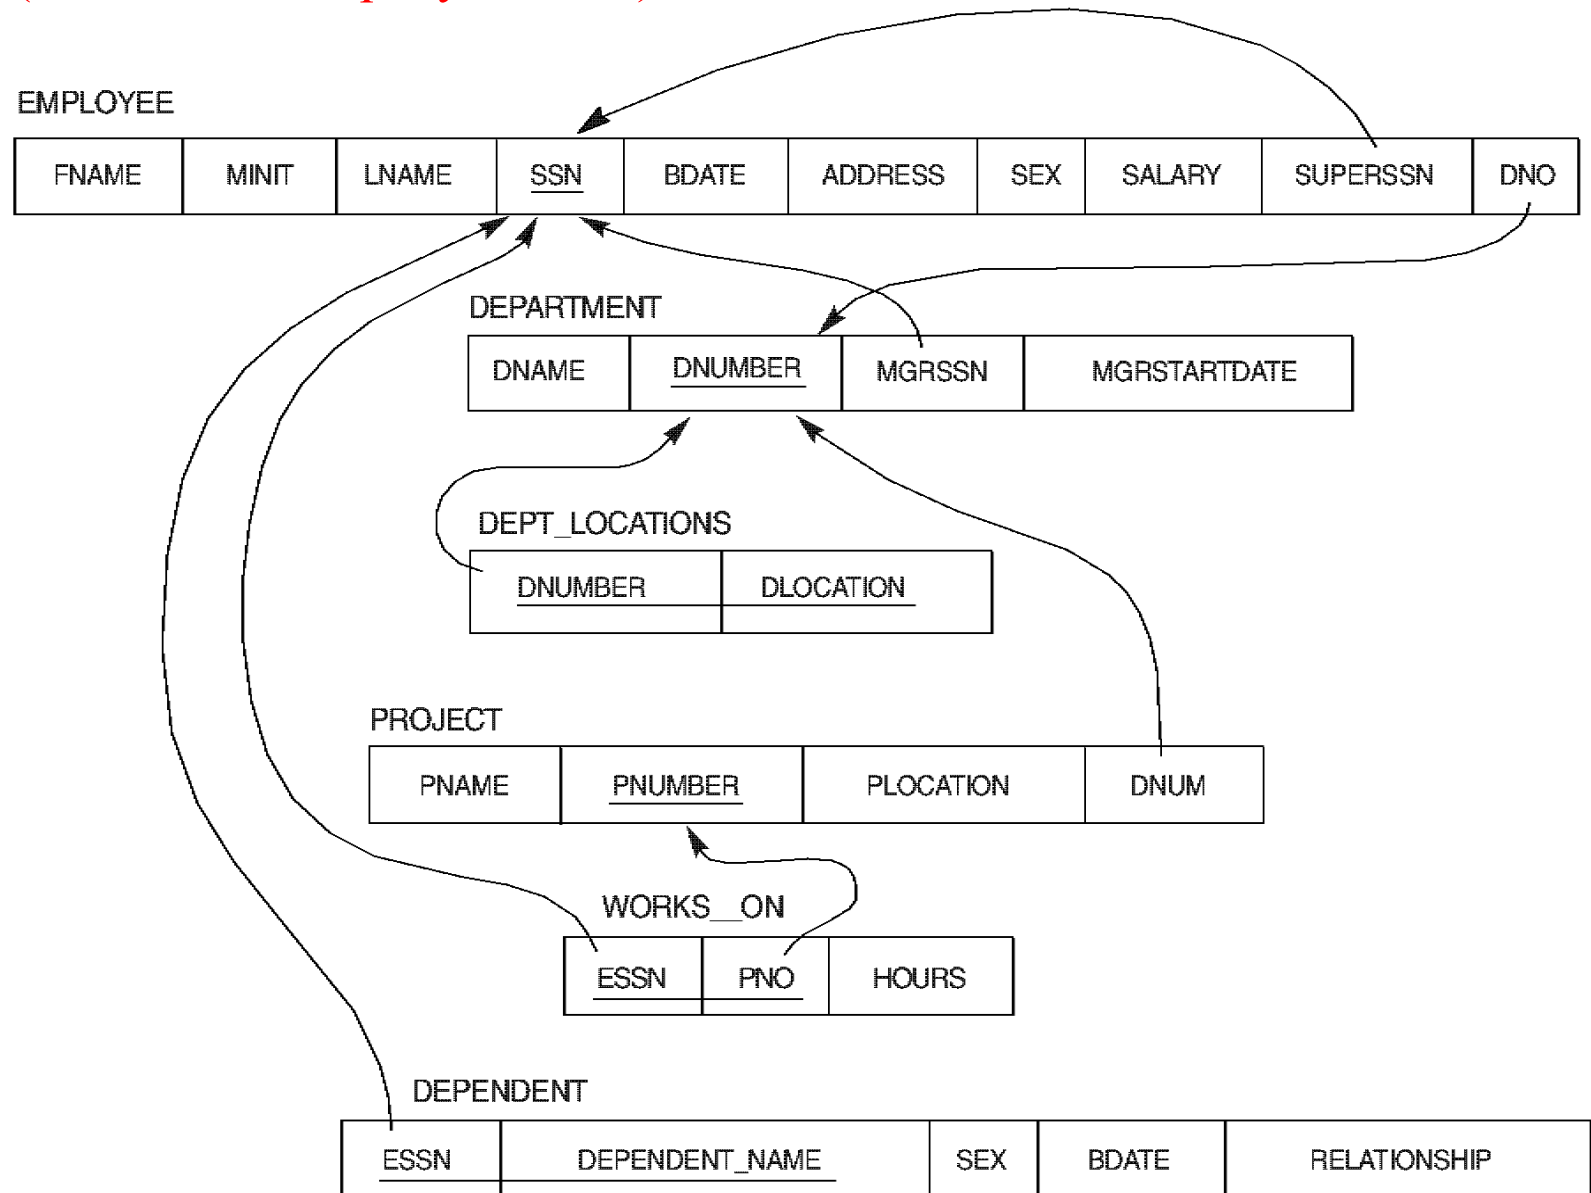
# Aggregate Functions (Cont.)

- Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

```
SELECT      MAX(SALARY), MIN(SALARY),  
            AVG(SALARY)  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNO=DNUMBER AND  
            DNAME='Research'
```

Note: the aggregate functions are applied to the relation consisting of all **pairs of tuples from Employee and Department** satisfying the condition in the WHERE clause

## (Reminder: company schema)



# Grouping (example)

## Employee

Name	Dept	Salary
Joe	Toys	45
Nick	PCs	50
Jim	Toys	35
Jack	PCs	40

*Find average salary of all employees*

```
SELECT Avg(Salary) AS AvgSal  
FROM Employee
```

AvgSal
--------

42.5
------

*Find the average salary for each department*

```
SELECT Dept, Avg(Salary) AS AvgSal  
FROM Employee  
GROUP BY Dept
```

Dept	AvgSal
Toys	40
PCs	45

# Grouping

- Allows to apply the aggregate functions *to subgroups of tuples in a relation*
- Each subgroup of tuples consists of the set of tuples that have *the same value* for the *grouping attribute(s)*
- The function is applied to each subgroup independently
- SQL has a **GROUP BY**-clause for specifying the grouping attributes
- Most systems require the grouping attributes to appear in the SELECT-clause
- Non-grouping attributes may only appear in the SELECT clause as arguments to aggregate functions

R	A	B	C
	a	b	0
	a	b	1
	a	b	2
	a	c	1
	a	c	3

select A, B, MAX(C) as M  
from R  
group by A , B

select A, MAX(C) as M  
from R  
group by A , B

select MAX(C) as M  
from R  
group by A , B

A	B	M
A	M	
M		

# Grouping (cont.)


- For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT DNO, COUNT (*) AS NUMEMP, AVG (SALARY) AS AVGSAL  
FROM      EMPLOYEE  
GROUP BY  DNO
```

- The EMPLOYEE tuples are divided into groups--each group having the same value for the grouping attribute DNO
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the aggregate functions to be applied on each group of tuples



## (Reminder: company schema)



# GROUPING Example

- For each project, retrieve the project number, project name, and the number of employees who work on that project.

```
SELECT      PNUMBER, PNAME, COUNT (*)  
FROM        PROJECT, WORKS_ON  
WHERE       PNUMBER=PNO  
GROUP BY    PNUMBER, PNAME
```

- Note: the grouping and functions are applied **on pairs of tuples from PROJECT, WORKS\_ON**

Subtlety: suppose PNO and ESSN do not form a key for WORKS\_ON  
Problem: will get duplicate employees

Works_on	ESSN	PNO	HOURS	PROJECT	PNAME, PNUMBER
	111-11-1111	001	20		Wiki 001
	111-11-1111	001	10		Geo 002
	22-22-2222	002	25		

Fix:

```
SELECT PNUMBER, PNAME, COUNT (DISTINCT ESSN)  
FROM PROJECT, WORKS_ON  
WHERE PNUMBER=PNO  
GROUP BY PNUMBER, PNAME
```

# THE HAVING-CLAUSE

- Sometimes we want to retrieve the values of aggregate functions for only those *groups that satisfy certain conditions*
- The HAVING-clause is used for specifying a selection condition on **groups** (rather than on individual tuples!)

# Aggregate Functions – Having Clause

- Find the name and average balance of all branches where the average account balance is more than \$1,200.

```
select      branch_name, avg (balance)  
from        account  
group by    branch_name  
HAVING      avg (balance) > 1200
```

Condition in HAVING clause use values of attributes in group-by clause and aggregate functions on the other attributes

# THE HAVING-CLAUSE (cont.)

- For each project *on which more than two employees work* , retrieve the project number, project name, and the number of employees who work on that project.

<b>SELECT</b>	<b>PNUMBER, PNAME, COUNT (*)</b>
<b>FROM</b>	<b>PROJECT, WORKS_ON</b>
<b>WHERE</b>	<b>PNUMBER=PNO</b>
<b>GROUP BY</b>	<b>PNUMBER, PNAME</b>
<b>HAVING</b>	<b>COUNT (*) &gt; 2</b>

Note: predicates in the having clause are applied after the formation of groups whereas predicates in the where clause are applied before forming groups

## Another example

*For each currently playing movie having more than 100 actors,  
find the number of theaters showing the movie*

```
SELECT m.Title, COUNT(distinct s.Theater) AS number  
FROM Schedule s, Movie m  
WHERE s.Title = m.Title  
GROUP BY m.Title  
HAVING COUNT(DISTINCT m.Actor) > 100
```

Aggregate is taken over pairs <s,m> with same Title

Schedule	Theater	Title
	Hillcrest	Star Wars
	Paloma	Star Wars

Movie	Title	Director	Actor
	Star Wars	Lucas	Ford
	Star Wars	Lucas	Fischer

**FROM** Schedule s, Movie m  
**WHERE** s.Title = m.Title

	Theater	Title	Director	Actor
	Hillcrest	Star Wars	Lucas	Ford
	Paloma	Star Wars	Lucas	Ford
	Hillcrest	Star Wars	Lucas	Fischer
	Paloma	Star Wars	Lucas	Fischer

**GROUP BY** m.Title

	Title	Theater	Director	Actor
	Star Wars	Hillcrest	Lucas	Ford
		Paloma	Lucas	Ford
		Hillcrest	Lucas	Fischer
		Paloma	Lucas	Fischer



# SQL Queries: Nesting

- The WHERE clause can contain predicates of the form
  - *attr/value IN <SQL query>*
  - *attr/value NOT IN <SQL query>*
- The IN predicate is satisfied if the *attr* or *value* appears in the result of the nested *<SQL query>*

## Examples:

*find directors of current movies*

```
SELECT director FROM Movie  
WHERE title IN
```

```
(SELECT title  
FROM schedule)
```

The nested query finds currently playing movies

# More examples

*Find actors playing in some movie by Bertolucci*

```
SELECT actor FROM Movie  
WHERE title IN
```

```
(SELECT title  
FROM Movie  
WHERE director = "Bertolucci")
```

The nested query finds the titles of movies by Bertolucci

In this case, can eliminate nesting:

```
SELECT actor FROM Movie  
WHERE title IN
```

```
(SELECT title  
FROM Movie  
WHERE director = "Bertolucci")
```

```
SELECT m1. actor  
FROM Movie m1, Movie m2  
WHERE m1.title = m2.title  
and m2.director = "Bertolucci"
```

Question: is nesting syntactic sugar? Can it always be eliminated?

A: yes B: no

Question: is nesting syntactic sugar? Can it always be eliminated?

A: yes B: no


Queries involving nesting but **no negation** can always be un-nested, unlike queries with nesting and negation

# Correlated nested queries

- If a condition in the WHERE-clause of a *nested query* references an attribute of a relation declared in the *outer query* , the two queries are said to be **correlated**
- The result of a correlated nested query **may be different** for each tuple (or combination of tuples) of the relation(s) the outer query
- E.g. DB Company: *Retrieve the name of each employee who has a dependent with the same first name as the employee.*

```
SELECT  E.FNAME, E.LNAME
FROM    EMPLOYEE E
WHERE   E.SSN IN
        (SELECT ESSN
         FROM   DEPENDENT
         WHERE  ESSN=E.SSN
          AND   E.FNAME=DEPENDENT_NAME)
```

## (Reminder: company schema)



## Correlated nested queries (cont.)

- Correlated queries using just the IN comparison operators  
*can still be unnested*
- For example, the previous query could be un-nested as follows:

```
SELECT  E.FNAME, E.LNAME  
FROM    EMPLOYEE E, DEPENDENT D  
WHERE   E.SSN=D.ESSN AND  
        E.FNAME=D.DEPENDENT_NAME
```

Use of NOT IN tests increases expressive power



## Simple use of NOT IN

*Find all movies in which Hitchcock does not act:*

```
SELECT title FROM Movie
```

```
Where title NOT IN
```

```
(SELECT title FROM Movie
```

```
WHERE actor = 'Hitchcock')
```

## Simple use of NOT IN

*Find all movies that are not currently playing:*

```
SELECT title FROM Movie
WHERE title NOT IN
      (SELECT title FROM Schedule)
```

# Why can't these be flattened?

Hand-waving “proof” :

1. Basic queries with no nesting are **monotonic**

the answer never decreases when the database increases  
 $DB1 \subseteq DB2$  implies  $Query(DB1) \subseteq Query(DB2)$

2. But queries using NOT IN are usually **not monotonic**

```
SELECT title FROM Movie
WHERE title NOT IN
      (SELECT title FROM Schedule)
```

If Schedule increases, the answer may decrease

# Recall semantics of basic queries:

```
SELECT  $a_1, \dots, a_n$   
FROM  $R_1, \dots, R_m$   
WHERE condition
```

for each tuple  $t_1$  in  $R_1$   
for each tuple  $t_2$  in  $R_2$

.....

for each tuple  $t_m$  in  $R_m$

if  $\text{condition}(t_1, t_2, \dots, t_m)$  then output in answer  
attributes  $a_1, \dots, a_n$  of  $t_1, \dots, t_m$

This is monotonic if condition has no nested queries

## Monotonic (A) or non-monotonic (B) ?

1. Find the theaters showing some movie by Fellini
2. Find the theaters showing only movies by Fellini
3. Find the actors who are also directors
4. Find the actors playing in some movie showing at Paloma
5. Find the actors playing in every movie by Bertolucci

Schedule	Theater	Title

Movie	Title	Director	Actor

## More complex use of NOT IN

*Find the names of employees with the maximum salary*

```
SELECT name FROM Employee
WHERE salary NOT IN
  (SELECT e.salary
   FROM Employee e, Employee f
   WHERE e.salary < f.salary)
```

Intuition: salary is maximum if it is **not** among salaries e.salary lower than some f.salary

## More complex use of **NOT IN**:

*Find actors playing in every movie by “Berto”*

SQL's way of saying this:

*find the actors for which  
there is no movie by  
Bertolucci in which  
they do not act*

OR equivalently:

*find the actors not among  
the actors for which there is  
some movie by Bertolucci  
in which they do not act*

```
SELECT Actor FROM Movie  
WHERE Actor NOT IN
```

```
(SELECT m1.Actor  
FROM Movie m1, Movie m2,  
WHERE m2.Director="Berto"  
AND m1.Actor NOT IN  
(SELECT Actor  
FROM Movie  
WHERE Title=m2.Title))
```

The shaded query finds actors for which there is some movie by “Berto” in which they do not act

The top lines complement the shaded part





- Another construct used with nesting: **EXISTS**

SELECT ...

FROM...

WHERE ... **EXISTS** (<query>)

Meaning of EXISTS:

**EXISTS** (<query>) is true iff the result of <query>  
is not empty.

**NOT EXISTS** (<query>) is true iff the result of <query>  
is empty.

Examples:

R	A	B
	1	1
	3	2
	5	7

EXISTS (select A from R where  $A < 4$ )

NOT EXISTS (select A from R where  $A > 10$ )

EXISTS (select \* from R where  $A < B$ )

NOT EXISTS (select \* from R where  $A < B$ )

**Example:** Find theaters showing a movie directed by Berto:

```
SELECT s.theater
FROM schedule s
WHERE EXISTS (SELECT * FROM movie
               WHERE movie.title = s.title AND
                     movie.director = 'Berto')
```

Schedule	Theater	Title

Movie	Title	Director	Actor

Example (Boolean predicate): Everybody likes Sara Lee



NOT EXISTS

```
(SELECT * FROM PERSON
WHERE NOT EXISTS
  (SELECT * FROM LIKES
   WHERE PERSON.name = LIKES.name
    AND brand = 'Sara Lee')
```

Example: Find the actors playing in every movie by Berto

```
SELECT a.actor FROM movie a
WHERE NOT EXISTS
  (SELECT * FROM movie m
   WHERE m.director = 'Berto' AND
        NOT EXISTS
          (SELECT *
           FROM movie t
           WHERE m.title = t.title
                AND t.actor = a.actor))
```

Movie	Title	Director	Actor

# Nested Queries: ANY and ALL

- $A \text{ op ANY } \langle \text{nested query} \rangle$  is satisfied if **there is** a value  $X$  in the result of the  $\langle \text{nested query} \rangle$  and the condition  $A \text{ op } X$  is satisfied
  - ANY aka SOME
- $A \text{ op ALL } \langle \text{nested query} \rangle$  is satisfied if **for every** value  $X$  in the result of the  $\langle \text{nested query} \rangle$  the condition  $A \text{ op } X$  is satisfied

*Find directors of currently playing movies*

```
SELECT Director
FROM Movie
WHERE Title = ANY
      SELECT Title
      FROM Schedule
```

*Find the employees with the highest salary*

```
SELECT Name
FROM Employee
WHERE Salary >= ALL
      SELECT Salary
      FROM Employee
```

# Nested Queries: Set Comparison

- *<nested query 1>* CONTAINS  
*<nested query 2>*

The original SQL as specified for SYSTEM R had a **CONTAINS** operator. This was dropped from the language, possibly because of the difficulty in implementing it efficiently

*Find actors playing in every movie  
by “Bertolucci”*

```
SELECT m1.Actor
FROM Movie m1
WHERE
    (SELECT Title
     FROM Movie
     WHERE Actor = m1.Actor)
CONTAINS
    (SELECT Title
     FROM Movie
     WHERE Director = “Berto”)
```

# Nested queries in FROM clause

SQL allows nested queries in the FROM clause

*Find directors of movies showing in Hillcrest:*

```
select m.director
from movie m,
(select title from schedule
where theater = 'Hillcrest') t
where m.title = t.title
```

Note: this is syntactic sugar and can be eliminated



# SQL: Union, Intersection, Difference

- *Union*
  - *<SQL query 1> UNION <SQL query 2>*
- *Intersection*
  - *<SQL query 1> INTERSECT <SQL query 2>*
- *Difference*
  - *<SQL query 1> EXCEPT <SQL query 2>*

These operations eliminate duplicates

*Find all actors or directors*

```
(SELECT Actor as Name  
FROM Movie)  
UNION  
(SELECT Director as Name  
FROM Movie)
```

*Find all actors who are not directors*

```
(SELECT Actor as Name  
FROM Movie)  
EXCEPT  
(SELECT Director as Name  
FROM Movie)
```

# SQL: Union, Intersection, Difference

- *Union*
  - *<SQL query 1> UNION ALL <SQL query 2>*
- *Intersection*
  - *<SQL query 1> INTERSECT ALL <SQL query 2>*
- *Difference*
  - *<SQL query 1> EXCEPT ALL <SQL query 2>*

To keep duplicates : ALL

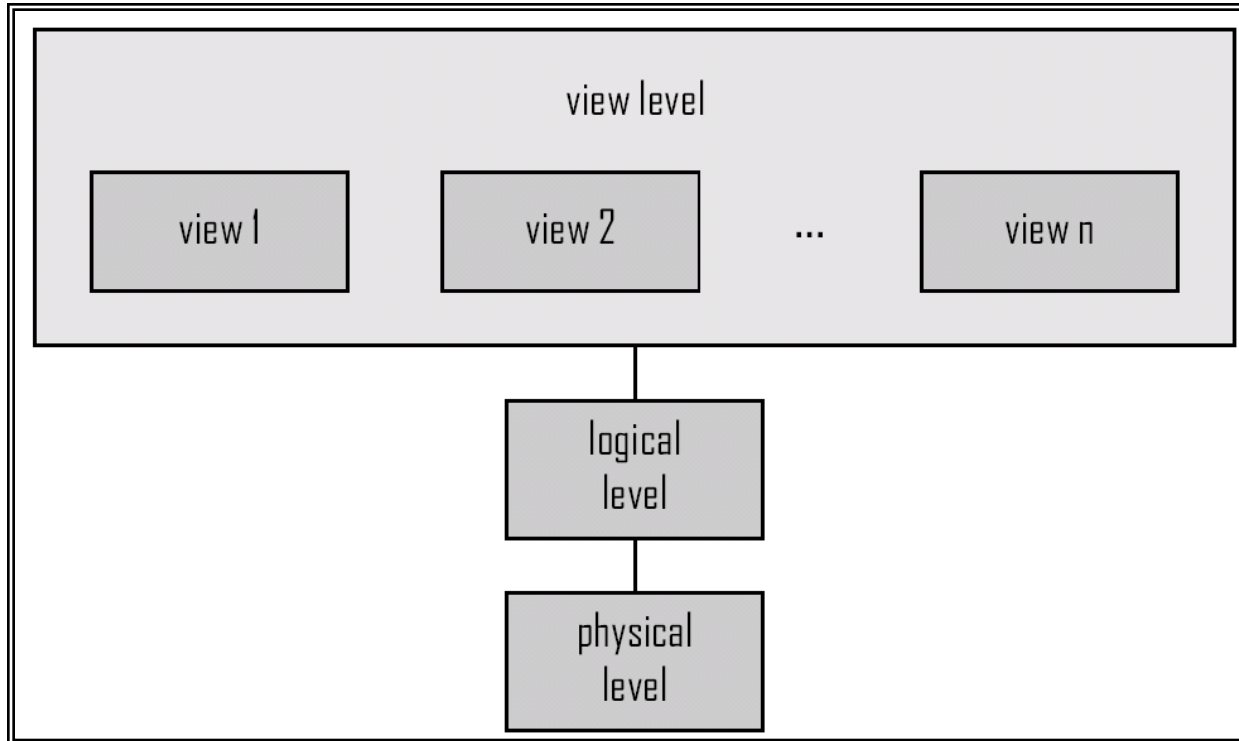
Example (union):

for each title in movie, find the number of theaters showing that title

schedule	theater title

movie	title director actor

# Basic Views (more later)



**Views** are a mechanism for customizing the database;  
also used for creating temporary virtual tables

# Views

- In some cases, it is not desirable for all users to see the entire logical model (i.e, all the actual relations stored in the database.)
- Consider a person who needs to know customers' loan numbers but has no need to see the loan amounts. This person should see a relation described, in SQL, by

```
(select customer_name, loan_number  
  from customer c, borrower b  
  where c.customer_id = b.customer_id)
```

- A **view** provides a mechanism to **hide or restructure** data for certain users.
- Any relation that is not in the database schema but is made visible to a user as a “virtual relation” is called a **view**.

# Bank relational schema

- $branch = (\underline{branch\_name}, branch\_city, assets)$
- $loan = (\underline{loan\_number}, branch\_name, amount)$
- $account = (\underline{account\_number}, branch\_name, balance)$
- $borrower = (\underline{customer\_id}, \underline{loan\_number})$
- $depositor = (\underline{customer\_id}, \underline{account\_number})$
- $customer = (customer\_id, customer\_name)$

# View Definition

- A view is defined using the **create view** statement which has the form

**create view  $V$  as < query expression >**

where  $V$  is the view name and <query expression> is any legal SQL query. A list of attribute names for  $V$  is optional.

- Once a view is defined, the view name can be used in queries
- Only limited updates can be applied to the view (more later)
- View definition is not the same as creating a new relation by evaluating the query expression: **the view contents is refreshed automatically when the database is updated**

# Examples

- A view consisting of bank branches and all their customers

**create view *all\_customers* as**

```
(select branch_name, customer_id
 from depositor d, account a
 where d.account_number = a.account_number)
union
(select branch_name, customer_id
 from borrower b, loan l
 where b.loan_number = l.loan_number)
```

- Find all customers of the La Jolla branch

```
select customer_id
from all_customers
where branch_name = 'La Jolla'
```



# Views can simplify complex queries

## Example

*find actors playing in every movie by “Berto”:*

```
SELECT Actor FROM Movie  
WHERE Actor NOT IN
```

```
(SELECT m1.Actor  
FROM Movie m1, Movie m2,  
WHERE m2.Director=“Berto”  
AND m1.Actor NOT IN  
  (SELECT Actor  
   FROM Movie  
   WHERE Title=m2.Title))
```

The shaded query finds actors NOT playing in some movie by “Berto”

# Same query using views

```
CREATE VIEW Bert-Movies AS  
SELECT title FROM Movie WHERE director = "Bertolucci" ;
```

```
CREATE VIEW Not-All-Berto AS  
SELECT m.actor FROM Movies m, Bert-Movies  
WHERE Bert-Movies.title NOT IN  
      (SELECT title FROM Movies  
       WHERE actor = m.actor);
```

```
SELECT actor FROM Movies WHERE actor NOT IN  
      (SELECT * FROM Not-All-Berto)
```

# Another syntax: the **with** clause

```
WITH Bertto-Movies AS
```

```
SELECT title FROM Movie WHERE director = “Bertoucci”
```

```
WITH Not-All-Berto AS
```

```
SELECT m.actor FROM Movies m, Bertto-Movies
```

```
WHERE Bertto-Movies.title NOT IN
```

```
    (SELECT title FROM Movies
```

```
        WHERE actor = m.actor)
```

```
SELECT actor FROM Movies WHERE actor NOT IN
```

```
    (SELECT * FROM Not-All-Berto) ;
```

Note: Bertto-Movies and Not-All-Berto are temporary tables, not views

Another example:

Employee	SSN	salary	dept
----------	-----	--------	------

Find the employees working in the departments with the highest average salary

```
create view Avg-sal as
select dept, AVG(salary) as average
from Employee
group by dept;
```

```
select e.SSN from Employee e, Avg-sal a
where e.dept = a.dept and a.average =
(select MAX(average) from Avg-sal)
```